

handwrittenDigit

October 25, 2015

```
In [12]: %pylab inline
# https://www.terminal.com/snapshot/86889ee10f164c1bebc47a37007e490d7a4c26fb20d6772c6eb232701
# pandas and numpy
# not so much of pandas but for read_csv which is more efficient than numpy.loadtxt
import numpy as np
import pandas as pd

# scikit-learn classifiers and cross validation utils
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.grid_search import GridSearchCV

# scikit-learn dimension reduction
from sklearn.decomposition import PCA

# scikit-learn dataset processing utils
from sklearn.preprocessing import MinMaxScaler
```

Populating the interactive namespace from numpy and matplotlib

```
In [6]: df = pd.read_csv('/root/machine-learning-classify-handwritten-digit/randomForest/Data/train.csv')
df = df.astype('float64')
df.shape
```

Out[6]: (42000, 785)

```
In [9]: df.ix[1000:1010,200:215]
```

```
Out[9]:
```

	pixel199	pixel200	pixel201	pixel202	pixel203	pixel204	pixel205	\
1000	0	0	0	0	0	0	0	
1001	0	0	0	0	0	0	0	
1002	0	0	0	0	0	0	0	
1003	0	0	0	0	0	0	0	
1004	0	0	0	0	0	0	0	
1005	0	0	0	46	221	253	253	
1006	0	0	0	0	0	0	0	
1007	0	0	0	0	0	0	0	
1008	0	0	0	0	0	0	0	
1009	0	0	0	0	0	0	241	
1010	0	0	0	0	0	0	0	

	pixel206	pixel207	pixel208	pixel209	pixel210	pixel211	pixel212	\
1000	0	0	0	211	239	42	0	
1001	0	0	0	0	0	0	0	


1002	0	0	0	51	254	254	139
1003	224	253	222	72	0	0	0
1004	0	2	132	254	215	55	0
1005	253	183	253	253	253	253	228
1006	0	12	171	252	253	252	252
1007	0	0	0	0	0	0	34
1008	0	0	0	10	29	41	141
1009	252	106	0	0	0	0	0
1010	0	0	0	0	0	0	57

	pixel213
1000	0
1001	0
1002	0
1003	0
1004	0
1005	0
1006	230
1007	82
1008	141
1009	253
1010	250

```
In [13]: f, (ax1, ax2, ax3, ax4) = subplots(ncols=4)
```

```
imsize = (28, 28)
ax1.matshow(np.reshape(df.ix[35023,1:], imsize), cmap='gray_r')
ax2.matshow(np.reshape(df.ix[1008,1:], imsize), cmap='gray_r')
ax3.matshow(np.reshape(df.ix[1009,1:], imsize), cmap='gray_r')
ax4.matshow(np.reshape(df.ix[1012,1:], imsize), cmap='gray_r')
```

```
f.tight_layout();
```



```
In [14]: #SVM
         #Create scaler and PCA objects
```

```
min_max_scaler = MinMaxScaler()
pca = PCA(n_components=80)
```

```
In [16]: X = min_max_scaler.fit_transform(df.ix[:9999,1:])
         X = pca.fit_transform(X)
```

```
y = df.ix[:9999,0]
```

```
In [17]: tuned_parameters = [{'kernel' : ['rbf'], 'gamma': [0.1, 1e-2, 1e-3], 'C': [10, 100, 1000]},
                             {'kernel' : ['poly'], 'degree' : [5, 9], 'C' : [1, 10]}]
```

```
svm = GridSearchCV( SVC(), tuned_parameters, cv=3, verbose=2 ).fit(X, y)
```

Fitting 3 folds for each of 13 candidates, totalling 39 fits

```
[CV] kernel=rbf, gamma=0.1, C=10 ...
[CV] ... kernel=rbf, gamma=0.1, C=10 - 19.2s
[CV] kernel=rbf, gamma=0.1, C=10 ...
[CV] ... kernel=rbf, gamma=0.1, C=10 - 19.3s
[CV] kernel=rbf, gamma=0.1, C=10 ...
[CV] ... kernel=rbf, gamma=0.1, C=10 - 19.7s
[CV] kernel=rbf, gamma=0.01, C=10 ...
[CV] ... kernel=rbf, gamma=0.01, C=10 - 3.4s
[CV] kernel=rbf, gamma=0.01, C=10 ...
[CV] ... kernel=rbf, gamma=0.01, C=10 - 3.4s
[CV] kernel=rbf, gamma=0.01, C=10 ...
[CV] ... kernel=rbf, gamma=0.01, C=10 - 3.5s
[CV] kernel=rbf, gamma=0.001, C=10 ...
[CV] ... kernel=rbf, gamma=0.001, C=10 - 3.5s
[CV] kernel=rbf, gamma=0.001, C=10 ...
[CV] ... kernel=rbf, gamma=0.001, C=10 - 3.6s
[CV] kernel=rbf, gamma=0.001, C=10 ...
[CV] ... kernel=rbf, gamma=0.001, C=10 - 3.6s
[CV] kernel=rbf, gamma=0.1, C=100 ...
[CV] ... kernel=rbf, gamma=0.1, C=100 - 19.2s
[CV] kernel=rbf, gamma=0.1, C=100 ...
[CV] ... kernel=rbf, gamma=0.1, C=100 - 19.3s
[CV] kernel=rbf, gamma=0.1, C=100 ...
[CV] ... kernel=rbf, gamma=0.1, C=100 - 19.5s
[CV] kernel=rbf, gamma=0.01, C=100 ...
[CV] ... kernel=rbf, gamma=0.01, C=100 - 3.4s
[CV] kernel=rbf, gamma=0.01, C=100 ...
[CV] ... kernel=rbf, gamma=0.01, C=100 - 3.4s
[CV] kernel=rbf, gamma=0.01, C=100 ...
[CV] ... kernel=rbf, gamma=0.01, C=100 - 3.5s
[CV] kernel=rbf, gamma=0.001, C=100 ...
[CV] ... kernel=rbf, gamma=0.001, C=100 - 2.9s
[CV] kernel=rbf, gamma=0.001, C=100 ...
[CV] ... kernel=rbf, gamma=0.001, C=100 - 2.9s
[CV] kernel=rbf, gamma=0.001, C=100 ...
[CV] ... kernel=rbf, gamma=0.001, C=100 - 2.9s
[CV] kernel=rbf, gamma=0.1, C=1000 ...
[CV] ... kernel=rbf, gamma=0.1, C=1000 - 19.2s
[CV] kernel=rbf, gamma=0.1, C=1000 ...
[CV] ... kernel=rbf, gamma=0.1, C=1000 - 19.2s
[CV] kernel=rbf, gamma=0.1, C=1000 ...
[CV] ... kernel=rbf, gamma=0.1, C=1000 - 19.2s
[CV] kernel=rbf, gamma=0.01, C=1000 ...
[CV] ... kernel=rbf, gamma=0.01, C=1000 - 3.4s
[CV] kernel=rbf, gamma=0.01, C=1000 ...
[CV] ... kernel=rbf, gamma=0.01, C=1000 - 3.4s
[CV] kernel=rbf, gamma=0.01, C=1000 ...
[CV] ... kernel=rbf, gamma=0.01, C=1000 - 3.4s
```

```

[CV] kernel=rbf, gamma=0.001, C=1000 ...
[CV] ... kernel=rbf, gamma=0.001, C=1000 - 3.0s
[CV] kernel=rbf, gamma=0.001, C=1000 ...
[CV] ... kernel=rbf, gamma=0.001, C=1000 - 3.0s
[CV] kernel=rbf, gamma=0.001, C=1000 ...
[CV] ... kernel=rbf, gamma=0.001, C=1000 - 3.0s
[CV] degree=5, kernel=poly, C=1 ...
[CV] ... degree=5, kernel=poly, C=1 - 14.1s
[CV] degree=5, kernel=poly, C=1 ...
[CV] ... degree=5, kernel=poly, C=1 - 14.3s
[CV] degree=5, kernel=poly, C=1 ...
[CV] ... degree=5, kernel=poly, C=1 - 14.0s
[CV] degree=9, kernel=poly, C=1 ...
[CV] ... degree=9, kernel=poly, C=1 - 16.2s
[CV] degree=9, kernel=poly, C=1 ...
[CV] ... degree=9, kernel=poly, C=1 - 16.6s
[CV] degree=9, kernel=poly, C=1 ...
[CV] ... degree=9, kernel=poly, C=1 - 16.6s
[CV] degree=5, kernel=poly, C=10 ...
[CV] ... degree=5, kernel=poly, C=10 - 9.9s
[CV] degree=5, kernel=poly, C=10 ...
[CV] ... degree=5, kernel=poly, C=10 - 10.0s
[CV] degree=5, kernel=poly, C=10 ...
[CV] ... degree=5, kernel=poly, C=10 - 9.9s
[CV] degree=9, kernel=poly, C=10 ...
[CV] ... degree=9, kernel=poly, C=10 - 16.1s
[CV] degree=9, kernel=poly, C=10 ...
[CV] ... degree=9, kernel=poly, C=10 - 16.0s
[CV] degree=9, kernel=poly, C=10 ...
[CV] ... degree=9, kernel=poly, C=10 - 15.9s

[Parallel(n_jobs=1)]: Done 1 jobs | elapsed: 19.2s
[Parallel(n_jobs=1)]: Done 39 out of 39 | elapsed: 6.7min finished

```

```
In [18]: svm.best_estimator_
```

```
Out[18]: SVC(C=100, cache_size=200, class_weight=None, coef0=0.0, degree=3, gamma=0.01,
            kernel='rbf', max_iter=-1, probability=False, random_state=None,
            shrinking=True, tol=0.001, verbose=False)
```

```
In [21]: Xt = min_max_scaler.transform(df.ix[35000:,1:])
        yt = df.ix[35000:,0]
```

```
        yp = svm.predict(pca.transform(Xt))
```

```
In [31]: print ('prediction accuracy: %.4f' % (1./len(yt) * sum( yp == yt )))
```

```
prediction accuracy: 0.9681
```


```
In [33]: print (yp[:30].astype(int))
```

```
[2 2 5 4 7 6 0 6 9 5 4 8 5 2 6 2 3 9 3 9 4 0 4 5 2 0 6 5 5 0]
```

```
In [35]: f, (ax1, ax2, ax3, ax4) = subplots(ncols=4)
```

```
    imsize = (28, 28)
    ax1.matshow(np.reshape(df.ix[35000,1:], imsize), cmap='gray_r')
    ax2.matshow(np.reshape(df.ix[35001,1:], imsize), cmap='gray_r')
    ax3.matshow(np.reshape(df.ix[35002,1:], imsize), cmap='gray_r')
    ax4.matshow(np.reshape(df.ix[35003,1:], imsize), cmap='gray_r')

    f.tight_layout();
```



```
In [37]: #RandomForest
```

```
min_max_scaler = MinMaxScaler()
pca = PCA(n_components=80)
X = min_max_scaler.fit_transform(df.ix[:9999,1:])
X = pca.fit_transform(X)

y = df.ix[:9999,0]

tuned_parameters = [{'max_features': ['sqrt', 'log2'], 'n_estimators': [1000, 1500]}]

rf = GridSearchCV(RandomForestClassifier(min_samples_split=1,
                                         n_jobs=-1),
                  tuned_parameters,
                  cv=3,
                  verbose=2).fit(X, y)
```

Fitting 3 folds for each of 4 candidates, totalling 12 fits

```
[CV] max_features=sqrt, n_estimators=1000 ...
[CV] ... max_features=sqrt, n_estimators=1000 - 18.7s
[CV] max_features=sqrt, n_estimators=1000 ...
[CV] ... max_features=sqrt, n_estimators=1000 - 19.3s
[CV] max_features=sqrt, n_estimators=1000 ...
[CV] ... max_features=sqrt, n_estimators=1000 - 18.8s
[CV] max_features=sqrt, n_estimators=1500 ...
[CV] ... max_features=sqrt, n_estimators=1500 - 28.2s
[CV] max_features=sqrt, n_estimators=1500 ...
[CV] ... max_features=sqrt, n_estimators=1500 - 28.0s
[CV] max_features=sqrt, n_estimators=1500 ...
[CV] ... max_features=sqrt, n_estimators=1500 - 29.4s
[CV] max_features=log2, n_estimators=1000 ...
[CV] ... max_features=log2, n_estimators=1000 - 15.7s
```

```
[CV] max_features=log2, n_estimators=1000 ...
[CV] ... max_features=log2, n_estimators=1000 - 16.1s
[CV] max_features=log2, n_estimators=1000 ...
[CV] ... max_features=log2, n_estimators=1000 - 16.1s
[CV] max_features=log2, n_estimators=1500 ...
[CV] ... max_features=log2, n_estimators=1500 - 24.7s
[CV] max_features=log2, n_estimators=1500 ...
[CV] ... max_features=log2, n_estimators=1500 - 24.7s
[CV] max_features=log2, n_estimators=1500 ...
[CV] ... max_features=log2, n_estimators=1500 - 24.6s

[Parallel(n_jobs=1)]: Done 1 jobs      | elapsed: 18.7s
[Parallel(n_jobs=1)]: Done 12 out of 12 | elapsed: 4.4min finished
```

```
In [39]: Xt = min_max_scaler.transform(df.ix[35000:,1:])
        yt = df.ix[35000:,0]

        yp = rf.predict(pca.transform(Xt))

        print ('prediction accuracy: %.4f' % (1./len(yt) * sum( yp == yt )))
```

prediction accuracy: 0.9420


```
In [40]: print (yt[:30].values.astype(int))

[2 2 5 4 7 6 0 6 4 5 4 8 5 2 6 2 3 9 3 9 4 0 4 5 2 0 6 5 5 0]
```

```
In [41]: f, (ax1, ax2, ax3, ax4) = subplots(ncols=4)

        imsize = (28, 28)
        ax1.matshow(np.reshape(df.ix[35000,1:], imsize), cmap='gray_r')
        ax2.matshow(np.reshape(df.ix[35001,1:], imsize), cmap='gray_r')
        ax3.matshow(np.reshape(df.ix[35002,1:], imsize), cmap='gray_r')
        ax4.matshow(np.reshape(df.ix[35003,1:], imsize), cmap='gray_r')

        f.tight_layout();
```



```
In [42]: #Submission attempts
        # First attempt: split data into three, build an SVM for each,
        # majority voting for end prediction --> Kaggle score 0.97429
```

```

scalers = []
pca_xfrms = []
clfs = []

# compute indexes where to split the data
ixs = np.arange(df.shape[0])
splits = np.split(ixs, [14000, 28000])

tuned_parameters = [{'gamma': [0.1, 1e-2, 1e-3], 'C': [10, 100, 1000]}]

# use all data for training
for s in splits:
    min_max_scaler = MinMaxScaler()
    pca = PCA(n_components=80)

    # get training subset
    X = df.ix[s,1:].copy()
    y = df.ix[s,0].copy()

    # all the transformations
    X = min_max_scaler.fit_transform(X)
    X = pca.fit_transform(X)

    # train the classifier
    svm = GridSearchCV( SVC(), tuned_parameters, cv=3, verbose=1 ).fit(X, y)

    # store scaler, PCA transformer, and SVM classifier for this subset
    scalers.append(min_max_scaler)
    pca_xfrms.append(pca)
    clfs.append(svm)

[Parallel(n_jobs=1)]: Done 1 jobs | elapsed: 37.1s
[Parallel(n_jobs=1)]: Done 27 out of 27 | elapsed: 7.3min finished

Fitting 3 folds for each of 9 candidates, totalling 27 fits
Fitting 3 folds for each of 9 candidates, totalling 27 fits

[Parallel(n_jobs=1)]: Done 1 jobs | elapsed: 37.1s
[Parallel(n_jobs=1)]: Done 27 out of 27 | elapsed: 7.3min finished

Fitting 3 folds for each of 9 candidates, totalling 27 fits

[Parallel(n_jobs=1)]: Done 1 jobs | elapsed: 36.9s
[Parallel(n_jobs=1)]: Done 27 out of 27 | elapsed: 7.2min finished

In [46]: df2 = pd.read_csv('/root/machine-learning-classify-handwritten-digit/randomForest/Data/test.csv')
df2 = df2.astype('float64')

In [47]: preds = np.zeros((len(clfs), df2.shape[0]))

i = 0
for scaler, xfrm, clf in zip(scalers, pca_xfrms, clfs):

```

```

Xt = df2.copy()
Xt = scaler.transform(Xt)
preds[i] = clf.predict(xfrm.transform(Xt))

i += 1

total_pred = [np.bincount(x).argmax() for x in preds.T.astype(int)]

In [49]: with open('/root/machine-learning-classify-handwritten-digit/randomForest/submission.csv', 'w') as f:
        f.write("ImageId,label\n")
        for x in enumerate(total_pred, start=1):
            f.write("%s,%s\n" % x)

In [ ]: # DeLong
        #      0.97371      2      Sun, 25 Oct 2015 16:40:42

```